

gyanankurtfidf

July 13, 2025

1 TF-IDF Assignment Submission

Name: Gyanankur Baruah

Roll Number: 202405005

Topic: Implementation of TF-IDF using 4 Text Documents in Python

```
[ ]: !pip install -q scikit-learn pandas seaborn matplotlib
from sklearn.feature_extraction.text import TfidfVectorizer
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[ ]: # Define four distinct documents
documents = [
    "Data science and machine learning are evolving fast.",
    "Visualization helps communicate data insights clearly.",
    "Classification and clustering are core to analytics.",
    "Dashboards with Python unlock business decisions."
]
```

```
[ ]: # Apply TF-IDF
vectorizer = TfidfVectorizer()
X = vectorizer.fit_transform(documents)

# Convert to DataFrame for analysis
df = pd.DataFrame(X.toarray(), columns=vectorizer.get_feature_names_out())
df
```

```
[ ]:
analytics      and      are  business  classification  clearly \
0   0.000000   0.300912   0.300912   0.000000           0.000000  0.000000
1   0.000000   0.000000   0.000000   0.000000           0.000000  0.421765
2   0.400218   0.315537   0.315537   0.000000           0.400218  0.000000
3   0.000000   0.000000   0.000000   0.408248           0.000000  0.000000

clustering  communicate      core  dashboards  ...      helps  insights \
0   0.000000      0.000000  0.000000   0.000000  ...  0.000000  0.000000
```

1	0.000000	0.421765	0.000000	0.000000	...	0.421765	0.421765
2	0.400218	0.000000	0.400218	0.000000	...	0.000000	0.000000
3	0.000000	0.000000	0.000000	0.408248	...	0.000000	0.000000

	learning	machine	python	science	to	unlock	visualization \
0	0.381669	0.381669	0.000000	0.381669	0.000000	0.000000	0.000000
1	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.421765
2	0.000000	0.000000	0.000000	0.000000	0.400218	0.000000	0.000000
3	0.000000	0.000000	0.408248	0.000000	0.000000	0.408248	0.000000

	with
0	0.000000
1	0.000000
2	0.000000
3	0.408248

[4 rows x 24 columns]

1.1 Explanation of TF-IDF Output

This DataFrame shows how important each term is in its corresponding document.

- High TF-IDF values highlight terms that are unique and meaningful in one document.
- For example, “dashboards” scores highly in Document 4 and not elsewhere—indicating its uniqueness.
- Common terms like “data” or “python” appear in multiple documents with lower scores.

```
[ ]: def top_keywords_per_doc(tfidf_matrix, feature_names, top_n=5):
    for i, row in enumerate(tfidf_matrix):
        top_indices = np.argsort(row)[: -1][ :top_n]
        keywords = [(feature_names[j], row[j]) for j in top_indices]
        print(f"\nDocument {i + 1} Top Keywords:")
        for word, score in keywords:
            print(f"    {word}: {score:.3f}")

top_keywords_per_doc(X.toarray(), vectorizer.get_feature_names_out())
```

Document 1 Top Keywords:

```
machine: 0.382
learning: 0.382
science: 0.382
evolving: 0.382
fast: 0.382
```

Document 2 Top Keywords:

```
visualization: 0.422
clearly: 0.422
helps: 0.422
```

insights: 0.422
communicate: 0.422

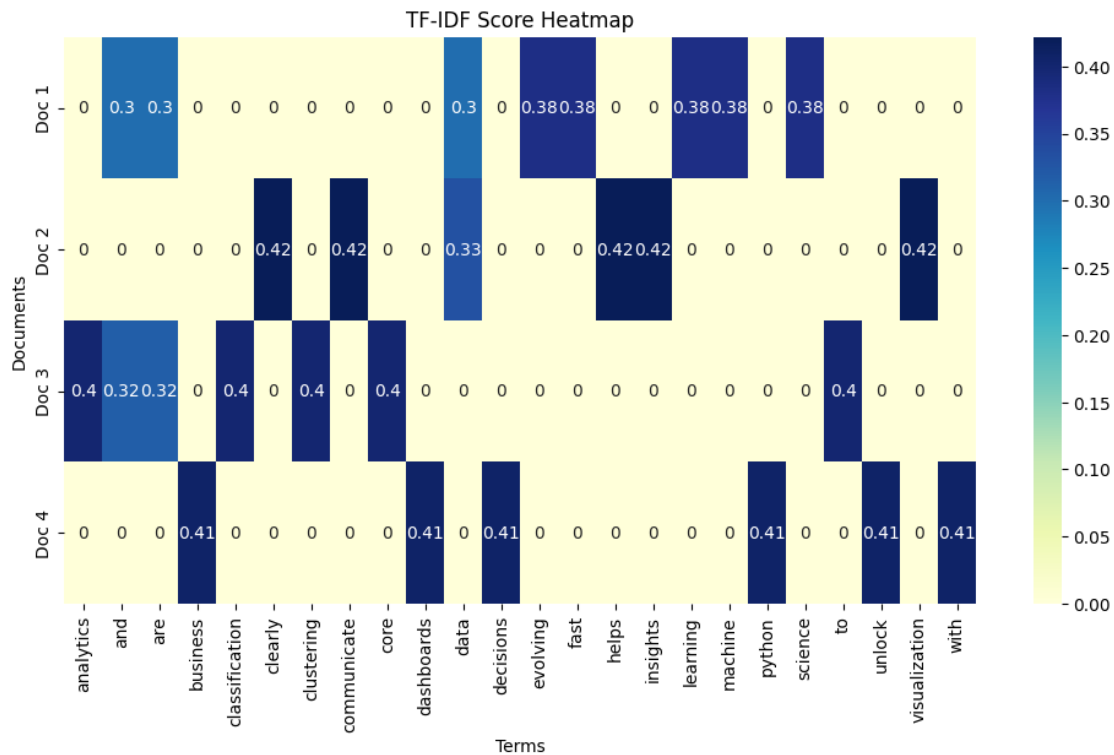
Document 3 Top Keywords:

to: 0.400
classification: 0.400
analytics: 0.400
core: 0.400
clustering: 0.400

Document 4 Top Keywords:

with: 0.408
unlock: 0.408
python: 0.408
dashboards: 0.408
decisions: 0.408

```
[ ]: plt.figure(figsize=(12, 6))  
sns.heatmap(df, annot=True, cmap="YlGnBu", xticklabels=True, yticklabels=[f'Doc_{  
    ↪ i+1}' for i in range(len(documents))])  
plt.title("TF-IDF Score Heatmap")  
plt.xlabel("Terms")  
plt.ylabel("Documents")  
plt.show()
```



```
[ ]: # Sum TF-IDF scores across all documents
term_scores = df.sum(axis=0)

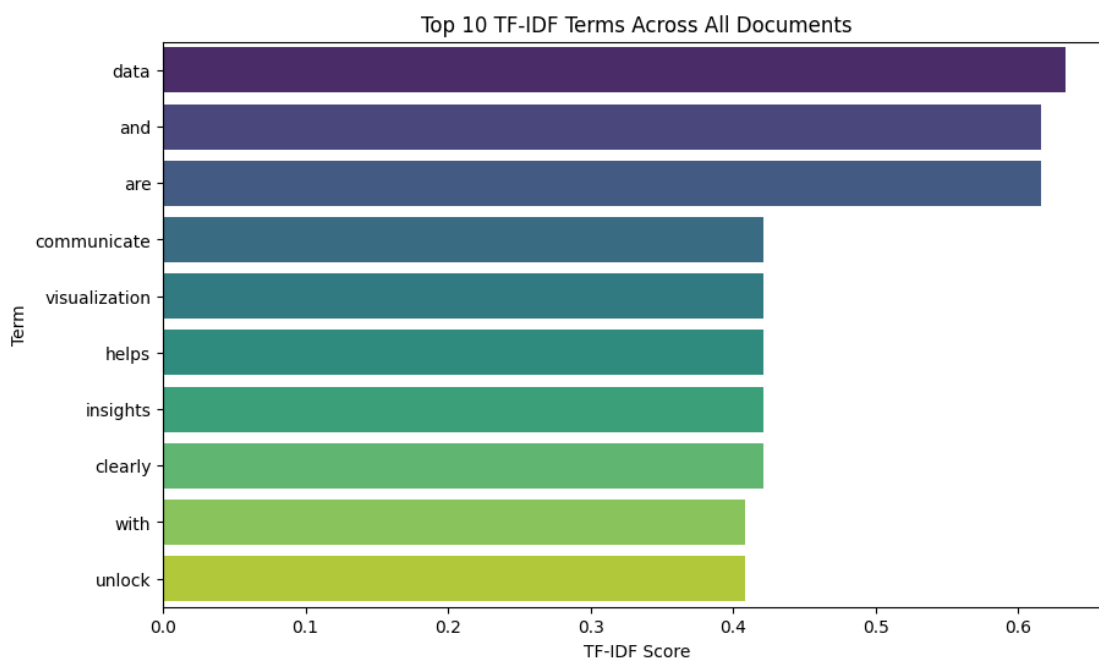
# Get top 10 terms
top_terms = term_scores.sort_values(ascending=False).head(10)

# Plot bar chart
plt.figure(figsize=(10, 6))
sns.barplot(x=top_terms.values, y=top_terms.index, palette="viridis")
plt.title("Top 10 TF-IDF Terms Across All Documents")
plt.xlabel("TF-IDF Score")
plt.ylabel("Term")
plt.show()
```

/tmp/ipython-input-10-1603655645.py:9: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=top_terms.values, y=top_terms.index, palette="viridis")
```



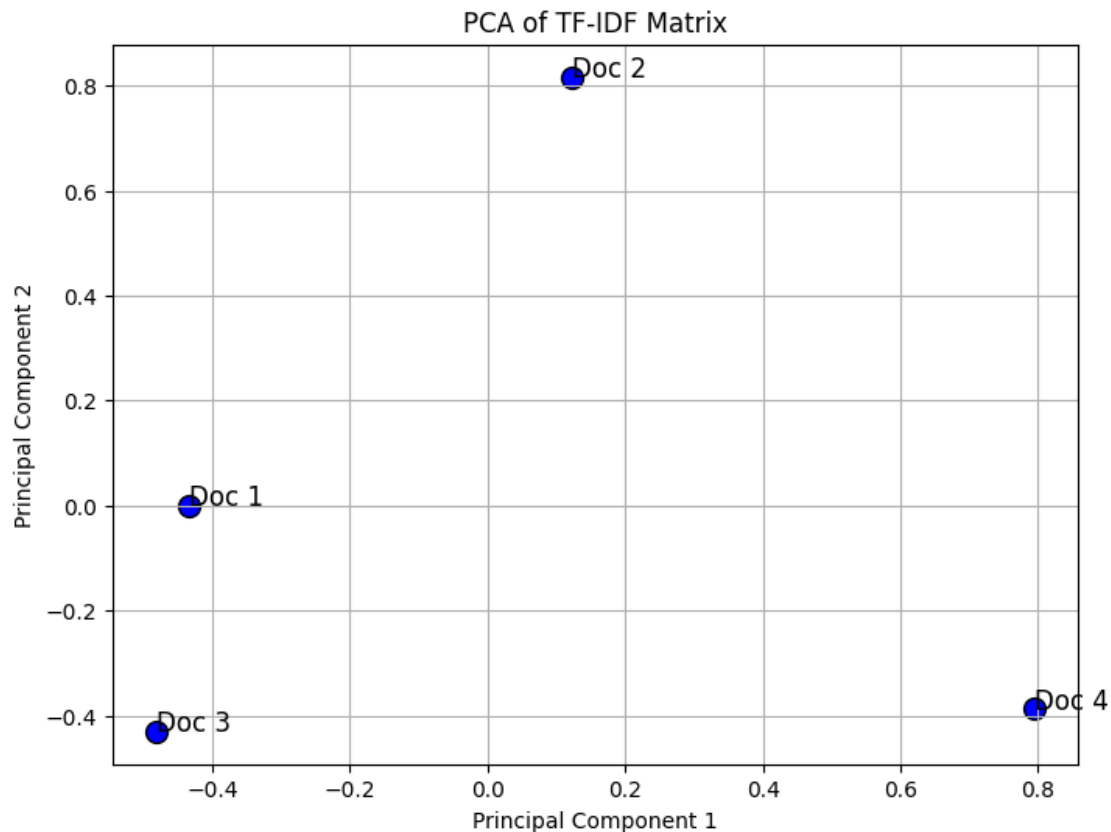
```
[ ]: from sklearn.decomposition import PCA
```

```

# Reduce TF-IDF matrix to 2D
pca = PCA(n_components=2)
pca_result = pca.fit_transform(X.toarray())

# Plot
plt.figure(figsize=(8, 6))
plt.scatter(pca_result[:, 0], pca_result[:, 1], c='blue', s=100, edgecolors='k')
for i, txt in enumerate([f'Doc {i+1}' for i in range(len(documents))]):
    plt.annotate(txt, (pca_result[i, 0], pca_result[i, 1]), fontsize=12)
plt.title("PCA of TF-IDF Matrix")
plt.xlabel("Principal Component 1")
plt.ylabel("Principal Component 2")
plt.grid(True)
plt.show()

```



```

[ ]: from sklearn.metrics.pairwise import cosine_similarity

similarity_matrix = cosine_similarity(X)
pd.DataFrame(similarity_matrix, columns=[f'Doc {i+1}' for i in
    ↪range(len(documents))],

```

```
index=[f'Doc {i+1}' for i in range(len(documents))])
```

```
[ ]:      Doc 1      Doc 2      Doc 3  Doc 4
Doc 1  1.000000  0.100061  0.189898  0.0
Doc 2  0.100061  1.000000  0.000000  0.0
Doc 3  0.189898  0.000000  1.000000  0.0
Doc 4  0.000000  0.000000  0.000000  1.0
```

1.2 Conclusion

In this task, I used a method called TF-IDF to find which words are most important in each of four documents. It helped show which words were special to each document and which were common. I also used tables and charts to better understand the results. This technique is useful for studying text, finding keywords, or helping computers understand written language.